

The AI Failure Tax

Silicon Valley profits from AI failure. China profits from the backlash.



By Andrew Horton

1 October 2026

Since 2022, Western capitals have treated artificial intelligence as a problem of containment. Convinced that geopolitical supremacy belongs to whoever controls the most advanced semiconductors and the largest models, Washington, Brussels and Canberra have weaponised export controls to deny China access to computing power. They are fighting the wrong war. The decisive vulnerability in the Western AI ecosystem is not that frontier models are too expensive, nor that export controls are too porous. It is that America's leading AI laboratories have built a commercial system that profits from its own shortcomings - a structural flaw that is quietly driving Western enterprise towards Chinese alternatives.

The architecture of Western AI rests on a simple premise: charge by the token. Every prompt, every response and every additional cycle of computation generates revenue. When large language models were little more than digital curiosities, billing by volume was a reasonable convention. Now that banks, governments, defence contractors and multinational corporations are deploying autonomous agents to execute business-critical workflows, token pricing has become a source of strategic distortion.

An increasing share of enterprise AI spending no longer purchases productive work. It purchases re-tries, hallucinations, verification loops and corrective actions. Organisations are not just paying for machine intelligence. They are paying for machine failure.

The scale of the problem is becoming impossible to ignore. On 13 July 2026, [McKinsey's QuantumBlack division](#) reported that roughly 60 per cent of the cost of an agentic task goes on checking outputs and repairing work the system failed to complete correctly on the first attempt. Some 93 per cent of organisations in [a McKinsey survey](#) were exceeding their AI budgets, and one in five respondents to its global survey had constrained their use of AI because of cost. [A separate audit](#) of nearly 10,000 customer service agent runs found that in the simpler domains, almost half of all failures involved systems falsely claiming to have completed tasks that remained undone. A company pays once for an agent to produce a defective outcome, then again to verify, identify and correct the error. Every hallucination increases revenue for the provider.

Call it the failure tax: the hidden levy that AI vendors extract from the operational friction of their own customers. The problem is economic. There is no evidence that any laboratory degrades its models deliberately, yet when a supplier's revenue grows with compute consumption, the penalty for unreliability falls on the buyer. The vendor captures the upside of activity, and the customer bears the cost of error. Capitalism rarely rewards such arrangements for long, and as chief information officers and boards grasp the true cost of the failure tax, their search for alternatives is creating a geopolitical consequence that export controls alone cannot prevent.

Enterprises are shifting predictable, high-volume workloads away from proprietary American platforms towards open-weight models they can run internally. A failed task still consumes electricity, but it no longer feeds a perpetual rent stream for a provider charging by consumption. That commercial refuge comes with an unexpected occupant, which is China. While America's AI champions built proprietary empires behind pay-per-token walls, Chinese developers flooded the global open-source ecosystem with increasingly capable models. In the year to spring 2026, they accounted for 41 per cent of all model downloads on [Hugging Face](#), the world's principal AI repository, and DeepSeek launched its V4 model at [roughly one-sixth the price](#) of leading American competitors.

The crucial question is no longer whether Chinese models are marginally better than American ones. The crucial question is whether they offer businesses an escape from Silicon Valley's pricing model. If they do, commercial gravity becomes a geopolitical force. [Governments across Asia](#) are already building national programs on cheaper Chinese foundations.

This is how influence spreads in the digital age. Not necessarily through superiority, but through dependency.

Chief financial officers are not conducting national security assessments. Procurement teams are not deciding the future balance of power between democratic and authoritarian systems. They are reducing costs and improving efficiency. Yet thousands of rational commercial decisions can collectively produce a strategic outcome that no policymaker intended. This is how influence spreads in the digital age. Not necessarily through superiority, but through dependency.

Many Western policymakers take comfort in sovereign hosting, the belief that a Chinese model running on domestic infrastructure somehow ceases to be Chinese. That is a category error. Model weights embody the assumptions, priorities and judgements accumulated during training. They are compressed cognitive frameworks, and every hospital, university, bank and defence contractor that embeds foreign models into its workflows adopts, to some degree, the intellectual architecture within them. Data exfiltration is the lesser danger; cognitive dependence is the greater. A nation does not need to control another country's data centres to shape its decision-making environment. It merely needs to provide the systems through which decisions are increasingly made.

History offers a precedent for escaping this trap. For decades, jet-engine manufacturers profited handsomely from maintenance and replacement parts, so an engine failure was a crisis for the airline and a revenue opportunity for the manufacturer. Rolls-Royce shattered that misalignment with Power by the Hour in 1962 and later [TotalCare](#), charging airlines a fixed rate for each hour an engine flies. Reliability became a commercial imperative, and every failure that once generated revenue became a direct hit to the manufacturer's balance sheet.

Artificial intelligence requires the same transformation. The future of enterprise AI cannot be built on charging customers for computational effort. It must be built on charging for verified outcomes. Vendors should be rewarded when systems complete tasks successfully—not when they consume additional tokens trying to recover from mistakes. Outcome-based contracts, service-level guarantees and automatic rebates for validated failures should become the standard. Some software companies are already moving: in April 2026, [Zendesk and HubSpot](#) switched their AI agents to charging per resolution. The frontier laboratories remain attached to the meter, and that attachment risks becoming a strategic error. Outcome-based contracts, service-level guarantees and automatic rebates for validated failures should become the standard. The smartest models in the world count for little if the business model surrounding them drives customers into the arms of a strategic competitor.



Governments should accelerate this transition. Public procurement agencies across NATO and the G7 should refuse to pay for algorithmic rework and require vendors to price according to successful task completion. Corporate boards should demand transparency on the true cost of AI outcomes rather than headline consumption figures. Any organisation incapable of quantifying how much of its AI expenditure is devoted to correcting machine error is managing blind.

A civilisation that permits its technology champions to monetise failure will eventually discover that customers seek alternatives. And when those alternatives arrive carrying the strategic interests of a rival power, semiconductor embargoes and export controls will offer little defence. The most consequential technology transfer of the coming decade may not occur because China steals Western innovation. It may occur because the West prices its own customers into China's arms.